

# From Images to Robot Coordinates: Dataset Design, Detector Selection and Monocular 3D Localisation for Automated EV Battery Disassembly

Yujie Wang, Komiljon Kosimov\*, Shiv Ashutosh Katiyar, and Quang Nguyen

Department of Mechanical Engineering, University of Birmingham, Birmingham B15 2TT, United Kingdom  
{yxw1516, kxk445, mxn498}@student.bham.ac.uk, s.a.katiyar@bham.ac.uk

\*Corresponding author

**Abstract**—Recycling electric vehicle (EV) batteries is still largely a manual job, which is slow, expensive and genuinely dangerous. Automating it starts with perception: a robot cannot unbolt or lift a battery module until it knows where that module is, in metres, in its own coordinate frame. This paper works through that problem end to end, and reports what actually mattered. First, we show that the accuracy figures usually quoted for battery detectors are optimistic: a detector scoring 0.818 mAP@50 on the single facility it was trained on drops to 0.277 on a 225-image benchmark spanning ten independent pack sources, a 66% relative loss. Second, we assemble a multi-source training set (4,425 images, 13 sources) and compare a conventional from-scratch detector against one built on a frozen self-supervised backbone. The transformer-based model reaches 0.502 mAP@50 against 0.410, but a per-class breakdown shows the two are indistinguishable on consistently annotated packs; the entire advantage is robustness when the annotation convention shifts. Third, we find that annotation convention, not image difficulty, dominates performance: per-source module accuracy spans 0.995 to 0.043, and the worst case is a labelling disagreement rather than a hard image. Finally, we close the gap from pixels to metres without a depth camera. A fiducial-referenced monocular transform recovers metric coordinates from a single RGB view and drives a UR5e manipulator. Across nine object-position configurations on real hardware, using only the RGB stream of an Intel RealSense D405, target-height MAE was 2.245 mm and stand-off MAE 1.556 mm, and the robot reached the intended target region in all nine cases. We also report a result that constrains how such systems should be validated: image-space reprojection error does not predict metric accuracy, with height errors of 5–7 mm occurring at reprojection residuals below 0.11 px. We release the annotation set (16,945 polygon masks) and all evaluation code.

**Index Terms**—robotic disassembly, electric vehicle batteries, circular economy, object detection, domain generalisation, monocular 3D localisation, fiducial markers, dataset quality

## I. Introduction

Electric vehicles are reaching end of life in large numbers, and their battery packs are both a hazard and an asset. A pack contains lithium, cobalt and nickel, all classified by the European Commission as critical raw materials with significant supply-chain risk [1], [2]. Recovering them matters economically and environmentally: recycling capacity demand in Europe is projected to reach 170 GWh by 2035 [3]. At the same time, EU Battery

Regulation 2023/1542 requires documented condition data across the battery lifecycle [4], which manual inspection struggles to deliver consistently.

Yet disassembly today is still mostly people with hand tools. It is slow, it does not scale, and it exposes workers to high voltage and the risk of thermal runaway. This is precisely the kind of task robots should be doing, and a substantial research effort is now directed at automating it [5], [7], [8].

Every such system starts with perception. Before a manipulator can unbolt, grasp or lift anything, it has to find the battery modules and the busbars that connect them. The standard approach is a supervised object detector trained on images collected at one facility or one experimental rig.

That practice hides a problem. A recycling line does not receive one pack type; it receives whatever arrives that day, from different manufacturers, in different states of disassembly, under whatever lighting the building has. A detector validated only on the distribution it was trained on tells you very little about the next pack through the door. Our own earlier single-facility system [13] had exactly this limitation, and this paper answers it.

A second problem follows immediately, and it is the one that matters for robotics. Even a perfect detector returns pixels, and a manipulator needs metres. That gap is normally bridged with an RGB-D camera, which adds cost and calibration work, and brings a failure mode of its own: the specular, dark and reflective surfaces inside a battery pack are exactly where consumer depth sensors degrade [20]. We show the bridge can instead be built from one RGB camera and a few planar markers placed once in the workcell.

We make four contributions:

- 1) We quantify how badly single-facility evaluation overstates accuracy for EV battery module detection: a 66% relative loss (Section V-A).
- 2) We compare training strategies and architectures on identical data, showing a frozen self-supervised backbone improves cross-facility accuracy and isolating why: robustness to annotation shift, not better detection of ordinary modules (Section V-B).

- 3) We show annotation convention is the dominant failure mode, with per-source accuracy spanning 0.995 to 0.043, and give the dataset audit that exposes it (Sections III-C and V-C).
- 4) We close the loop to robot coordinates with a monocular, fiducial-referenced transform needing no depth sensor, validated on a UR5e across nine object-position configurations (Section V-D).

Relation to our prior work. Our earlier paper [13] reported a CPU-deployable two-stage pipeline for a single facility, including condition grading, which we deliberately do not revisit here. This paper instead addresses the limitation that work could not test, cross-facility transfer, and adds the metric localisation stage needed to drive a manipulator. The dataset, benchmark, architectures and localisation method are all new.

## II. Background and Related Work

### A. Automating battery disassembly

Zang et al. [5] survey robotic EV battery disassembly and conclude that the process remains predominantly manual, with perception and task uncertainty as the principal obstacles. Wegener et al. [6] established early on that pack designs vary enough between manufacturers that fixed automation does not transfer, which is the practical reason perception must generalise. Tan et al. [7] demonstrate a collaborative robot that unfastens and collects screws by fusing vision with force/torque sensing, reporting over 90% unscrewing efficiency. RAPID [8] integrates open-vocabulary detection with RGB-D sensing on a gantry manipulator and reports 0.9757 mAP@50 for component detection.

RAPID’s manipulation numbers set the honest scope for any perception claim: at that detection accuracy, one-shot vision-driven grasping still succeeded only 57% of the time, rising to 83% with visual servoing. Perception accuracy is necessary but plainly not sufficient, so we restrict our claims to perception and localisation.

### B. Perception for battery components

Most published work targets fasteners rather than modules. Screw detection has been addressed with Faster R-CNN plus rotation edge similarity [9], with YOLOX-family models [10], and in a recent YOLOv11 versus Faster R-CNN benchmark [11]. Module- and busbar-level perception is comparatively under-explored; the closest work applies instance segmentation with point-cloud registration to grasp busbars from a single module type [12].

A recurring theme is data scarcity. Every study above collected and hand-annotated its own images – 1,719 in [10] – and none released them. [11] resorts to laptop screws as a proxy domain and states plainly that the field is hampered by a lack of public datasets. That motivates our release.

TABLE I  
Principal data sources (13 total, ten pack families)

| Source                  | Imgs  | Pack family        |
|-------------------------|-------|--------------------|
| ev-battery-comp-detect. | 1,045 | multi-manufacturer |
| ev-battery-components   | 680   | Audi/BMW/Chevy     |
| last_exp4               | 483   | multi-manufacturer |
| ev-battery-pack         | 435   | single lab rig     |
| battery_comp            | 280   | Nissan Leaf        |
| ue_d1_defect            | 237   | module close-ups   |
| final_mobilenet         | 216   | multi-manufacturer |
| automated-disassembly   | 117   | disassembly cell   |
| bmw_i3, ue_rav4         | 207   | BMW i3, RAV4       |
| Others (3 sources)      | 188   | mixed              |

### C. Generalisation, and why data may matter more than architecture

Domain shift is a well-documented failure mode [14], [15]: models that appear strong on their own test split degrade sharply on data collected differently. Self-supervised backbones such as DINOv2 [17] learn features that transfer without task-specific fine-tuning, and detection transformers built on them [18] are reported to generalise better than from-scratch detectors in low-data regimes. Whether that holds on the cluttered, specular imagery of battery internals is what we test.

Running alongside this is the data-centric argument [16] that label quality limits performance more than model capacity. Our own prior work found the same effect [13]; here we quantify it across facilities.

### D. Recovering metric pose from images

Vision-guided disassembly usually obtains depth from RGB-D [8] or from point-cloud registration against a reference model [12]. Both add hardware and calibration cost, and structured-light and time-of-flight sensors are unreliable on the specular metal and dark polymer that dominate battery internals [20]. Planar fiducials such as ArUco [21] offer an alternative: a marker of known geometry fixes a metric frame from a single view via perspective- $n$ -point [22], and is widely used for camera pose estimation and hand-eye calibration. We take this route and quantify what it costs.

## III. Dataset: Construction and Analysis

### A. Sources

We assembled 4,425 training images from 13 public sources spanning ten EV pack families (Table I), annotated with two classes, module and busbar. Licences differ across sources and one is non-commercial; we record per-source provenance so any single source can be excluded.

### B. A cross-facility test set, and why deduplication is not optional

Reported accuracy only means something against a benchmark that resembles deployment. We therefore built

a cross-facility test set of 225 images (780 instances) drawn from all sources and held out entirely from training. We verified with perceptual hashing (dHash, Hamming distance  $\leq 6$ ) that no near-duplicate crosses the train/test boundary for any source we report per-source results on.

This proved necessary rather than precautionary. One candidate source, a 712-image public collection covering 19 vehicle types, re-publishes images from another source already in our set: 78 were byte-identical duplicates of our training images. Merging without checking would have leaked training data straight into the test set.

### C. What a dataset audit reveals

We audited all 4,425 training images for label and image quality. Of these, 23.5% carry no labels, 12.3% score low on a blur metric, 20.4% are greyscale, and none contain degenerate boxes. Two findings changed how we treated the data.

First, blur and greyscale are largely false alarms. The lowest-scoring images are smooth, closed pack lids – a Laplacian variance metric reads flat metal as “blurry” – and the greyscale images come from two monochrome sources. Both add useful variation; deleting them on the metric alone would be a mistake.

Second, the harmful category is unlabelled images that nonetheless contain visible components. Some of the 23.5% is legitimate background, which is healthy in moderation because it suppresses false positives. But inspection confirmed open packs with clearly visible busbar assemblies carrying no annotations at all, frames that actively teach the detector that busbars are background and plausibly contribute to the weak busbar scores below.

We also found 16,945 of 19,522 annotations are polygon masks (5–151 vertices) rather than boxes, a richer signal than detection consumes. We convert them to enclosing boxes for evaluation, and note that an early version of our own pipeline silently discarded every polygon label, trained on 20% of the annotations, and produced an inflated score that looked plausible. The failure is easy to introduce and hard to notice.

## IV. Method

### A. Pipeline

The system has two stages. A detector localises module and busbar instances in the image. A localisation stage then lifts each detection into metric robot-frame coordinates, so the output is directly usable by a manipulator.

### B. Detectors and training strategies compared

Single-facility specialist. YOLOv8n trained only on the single-facility source, following the conventional approach and our own prior work [13].

Multi-source YOLO11n. YOLO11n (2.6 M parameters) [19] trained from COCO initialisation on all 13 sources, 150 epochs, 640px, SGD.

Multi-source RF-DETR. RF-DETR-Nano [18], a real-time detection transformer with a frozen DINOv2 [17] backbone, fine-tuned on the same data at 384px for 60 epochs.

### C. Evaluation protocol

Confidence scores are not comparable across architectures, so every model is scored with a single metric implementation at a common confidence threshold of 0.05, on identical images, against identical ground truth including polygon-derived boxes. Mixing metric implementations across architectures silently produced a spurious result in our own early experiments.

### D. Monocular fiducial-referenced localisation

Let  $\mathbf{K}$  be the intrinsic matrix of a calibrated monocular camera. We place  $N$  ArUco markers [21] of known side length at known positions on the workcell plane. Detecting their corners gives  $4N$  correspondences between image points  $\mathbf{u}_i$  and world points  $\mathbf{X}_i$ , from which the camera pose  $(\mathbf{R}, \mathbf{t})$  follows by minimising reprojection error [22],

$$(\mathbf{R}^*, \mathbf{t}^*) = \arg \min_{\mathbf{R}, \mathbf{t}} \sum_{i=1}^{4N} \|\mathbf{u}_i - \pi(\mathbf{K}, \mathbf{R}, \mathbf{t}, \mathbf{X}_i)\|^2, \quad (1)$$

where  $\pi(\cdot)$  is the perspective projection. Using all markers jointly rather than one at a time over-constrains the fit and averages out per-marker corner noise.

Four markers (ID 0–3, 80 mm, DICT\_4X4\_50) fix the Board frame; the target pose is solved separately from a 40 mm marker using an IPPE planar-PnP solution, since a square planar target admits competing solutions under perspective projection. Candidates are screened on reprojection error, consistency of the target normal with the Board normal, and continuity of estimated height. A target position in the Board frame follows as

$$\mathbf{p}_B = \mathbf{R}_{CB}^\top (\mathbf{t}_{CT} - \mathbf{t}_{CB}), \quad (2)$$

with  $\mathbf{R}_{CB}, \mathbf{t}_{CB}$  the Board-to-camera pose and  $\mathbf{t}_{CT}$  the target translation in camera coordinates. A frame is accepted only when both Board and target reprojection RMSE are  $\leq 2.0$  px, and the coordinate-wise median of 30 accepted samples forms the commanded target, which suppresses isolated corner-detection outliers.

### E. Transformation to the robot frame and motion

Board coordinates are mapped to the UR5e Base frame by a calibrated planar rigid transformation, appropriate here because camera, board and robot base are fixed and motion is confined to one planar work surface with fixed tool orientation. The vertical coordinate is handled separately: the Board plane sits at Base  $Z = 82.04$  mm, and the commanded height is that plane plus the visually estimated target height plus a nominal 50 mm stand-off. Separating the planar and vertical paths lets their errors be attributed independently, which proves important in Section V-D. Motion runs in two bounded stages, an

TABLE II  
Cross-facility generalisation gap

| Model                      | In-domain<br>mAP@50 | Cross-facility<br>mAP@50 |
|----------------------------|---------------------|--------------------------|
| Single-facility specialist | 0.818               | 0.277                    |
| Multi-source YOLO11n       | 0.195               | 0.410                    |
| Multi-source RF-DETR       | 0.541               | 0.502                    |

TABLE III  
Detector comparison, cross-facility test (225 images)

| Model               | mAP<br>@50 | mAP<br>@50–95 | mod.  | bus.  |
|---------------------|------------|---------------|-------|-------|
| Specialist          | 0.277      | –             | 0.231 | –     |
| YOLO11n (multi-src) | 0.410      | 0.256         | 0.547 | 0.304 |
| RF-DETR (DINOv2)    | 0.502      | 0.312         | 0.680 | 0.353 |

approach above the target at 0.020 m/s followed by a vertical descent at 0.005 m/s, with workspace limits enforced before execution.

## V. Results

### A. How much single-facility evaluation overstates accuracy

Table II quantifies the central problem. The specialist scores 0.818 mAP@50 on the single-facility benchmark it was trained and validated on, and 0.277 on the cross-facility benchmark: a 66% relative loss. Its in-domain number simply does not predict deployment behaviour.

The reverse also holds, and is worth stating: the multi-source models score worse in-domain than the specialist. Training for breadth costs peak accuracy on any one pack, and which trade you want depends on whether your line sees one pack type or many. The size of that drop for YOLO11n (0.195) is larger than breadth alone explains, and is not a training failure: the in-domain benchmark happens to be the one source whose annotation convention diverges from the rest, a point we return to in Section V-C. RF-DETR, evaluated on exactly the same images, retains 0.541.

### B. Comparing architectures

Under the identical protocol above, RF-DETR reaches 0.502 mAP@50 against 0.410 for YOLO11n (Table III), and leads on the stricter mAP@50–95 (0.312 vs 0.256).

Taken alone that comparison is misleading. Restricting evaluation to the seven sources sharing a consistent annotation convention (177 images), the two models are effectively tied: overall 0.558 (RF-DETR) vs 0.544 (YOLO11n), module 0.771 vs 0.774, busbar 0.366 vs 0.344. The DINOv2 backbone confers no advantage on conventionally annotated packs.

The whole aggregate advantage therefore comes from behaviour under annotation shift. Adding the outlier source drops YOLO11n’s module score from 0.774 to

0.547, but RF-DETR’s only from 0.771 to 0.680. The claim we can defend is narrow and specific: a frozen self-supervised backbone does not detect ordinary modules better, but it degrades far more gracefully when the target definition moves. On a line receiving heterogeneous packs, graceful degradation is the property worth paying for.

### C. Annotation convention as the dominant failure mode

Per-source module mAP@50 is excellent on most packs – 0.995, 0.910, 0.873, 0.749 – and collapses to 0.043 on exactly one source.

That outlier is not a harder image. It annotates small internal sub-components in extreme macro views as “modules”, while every other source annotates pack-level modules. The detector learned the majority convention and applied it consistently; the collapse is a disagreement about the target, not a failure of perception.

Two consequences follow. Aggregate mAP over sources with inconsistent conventions is not meaningful, so per-source reporting is essential. And an attempt to repair the outlier automatically with an open-vocabulary grounding model failed outright: on greyscale macro imagery the proposals were unusable. Convention conflicts are not cheaply fixable after the fact, which argues for settling annotation guidelines before collection rather than after.

Busbar is the weaker class for every model (0.30–0.37), and 61% of busbar instances in the test set come from one source, with only four of eight sources annotating busbars at all – the same coverage problem in a different guise.

### D. From detection to metric robot coordinates

A detection is only actionable once expressed in the robot’s frame. We evaluate the localisation stage of Section IV-D on the hardware shown in Fig. 2: an Intel RealSense D405 fixed above the work surface, a UR5e manipulator, and four reference markers bounding the workspace. Only the RGB stream is used, at  $1280 \times 720$  and 30 fps; the depth stream is deliberately left unused so that what is measured is monocular fiducial geometry rather than sensed depth. Board calibration over 20 views converged to a reprojection RMSE of 0.248 px.

Three objects, two with regular top surfaces (56 and 63 mm) and one irregular, were each placed at three workspace positions, giving nine object-position configurations, 270 accepted visual samples and nine physical robot executions.

Visual repeatability. Within-configuration standard deviations were 0.103, 0.340 and 0.505 mm in  $X$ ,  $Y$  and  $Z$ . Mean Board and target reprojection RMSE were 0.388 and 0.191 px, all within the 2.0 px acceptance threshold. Short-term estimates are therefore stable under fixed-camera conditions.

Accuracy and execution. Table IV and Fig. 1 summarise. Target-height MAE was 2.245 mm (RMSE 3.177, signed  $-0.157 \pm 3.365$  mm), stand-off MAE 1.556 mm against the

TABLE IV  
Localisation and execution on hardware (9 configurations)

| Metric                        | Result               |
|-------------------------------|----------------------|
| Accepted visual samples       | 270 (30 per config.) |
| Height MAE / RMSE             | 2.245 / 3.177 mm     |
| Planar error, autonomous (O1) | 3.605 mm (max 6.091) |
| Planar calibration RMSE       | 5.877 mm             |
| Stand-off MAE / RMSE          | 1.556 / 1.944 mm     |
| Target reaches                | 9/9 (100%)           |




Fig. 1. Localisation on hardware. (a) Signed height error for all nine object-position configurations; O2 dominates the spread. (b) Error budget. The planar figure is the autonomous O1 subset, not the post-alignment measurements.

nominal 50 mm, and the robot reached the intended target region in all nine configurations.

Reprojection error does not predict metric accuracy. This is the result we consider most consequential for how such systems are validated. Two configurations produced height errors of  $-5.330$  and  $+6.910$  mm at target reprojection residuals of just 0.101 and 0.052 px. Planar PnP depth is sensitive to corner geometry, target tilt, marker flatness and the physical definition of the target plane, none of which a low image-space residual detects. A system validated on reprojection error alone may be substantially worse in robot space than it appears.

On the planar figure. Six of the nine planar measurements followed manual flange-centre alignment after automatic arrival, so they cannot be read as autonomous accuracy; we base our claim on the three genuinely autonomous measurements, which average 3.605 mm with a 6.091 mm maximum against a 5.877 mm calibration RMSE. The defensible statement is that the system moves the robot reliably into the target neighbourhood, not that it achieves sub-millimetre absolute accuracy.

The error budget points somewhere useful. Board and target reprojection residuals averaged 0.388 and 0.191 px, comfortably inside tolerance, yet the planar calibration fit carried an RMSE of 5.877 mm, larger than any error we measured downstream of it. The reference frame, not the vision, dominates. Improving marker placement and calibration is therefore a more productive target than sharpening the detector – not obvious in advance, and easy to get backwards. Set against the manipulation literature, RAPID reports 0.9757 mAP@50 yet 57% one-shot grasp success [8]: the loss occurs between perception




Fig. 2. Vision-guided workcell. An Intel RealSense D405 observes the surface from above; four reference markers bound the workspace and define the Board frame. Only the RGB stream is used. The UR5e approaches the target from above under a bounded two-stage motion.

and actuation. A 2.245 mm height MAE with a 4 mm worst-case stand-off sits inside the closing tolerance of a compliant gripper at this scale, so the monocular route is viable for coarse positioning with final alignment handed to force feedback or visual servoing.

#### E. Deployment and negative results

INT8 quantisation of the detector to ONNX cut latency 20% and size 46% (44.8 ms, 3.4 MB on CPU) for a loss of 0.007 mAP@50. Localisation adds only marker detection and a PnP solve, negligible against detector inference.

Three failures constrain the design space. Open-vocabulary detection is unusable here: YOLO-World, prompted with “ev battery module” and “busbar”, scored 0.004 mAP@50, battery internals being too far out of distribution. Naive multi-source merging reduced module mAP@50 to 0.331, below the single-source baseline: more data made the model worse. And copy-paste and glare augmentation reduced mAP here, despite being widely reported as beneficial elsewhere.

#### VI. Limitations

Our claims concern perception and localisation. We report metric accuracy and target reaching, not grasp success: the manipulator approached each target but performed no grasp, and gripper compliance and contact dynamics are outside our scope.

The localisation experiments cover nine configurations with one robot execution each: enough to demonstrate the chain end to end, too few to characterise workspace-wide error statistically. Planar and stand-off ground truth

were measured manually, so these are experimental validation values rather than traceable metrology; an external tracker would be needed to separate calibration error from robot repeatability and operator measurement error. We validated approach and arrival, not grasping, so task-level disassembly tolerance cannot be inferred from stand-off accuracy alone.

The localisation stage used a fiducial on the target object, whereas the detector operates markerless. Substituting a detected mask centroid for that marker is the obvious next step, and we expect it to cost accuracy: marker corners resolve to a fraction of a pixel, a learned centroid does not. The method also assumes targets rest on a support of known height, which stacked or tilted packs violate.

Finally, our data is assembled from public sources rather than collected under controlled conditions, so pack-type coverage is uneven and licences vary.

## VII. Conclusion

Three conclusions follow. Single-facility evaluation overstates deployment accuracy badly enough to mislead: the same detector lost two thirds of its accuracy when moved across pack sources. A frozen self-supervised backbone helps, but not in the way one might assume, buying graceful degradation under annotation shift rather than better detection of ordinary modules. And annotation convention, not architecture, proved the dominant failure mode, which points at dataset design as the first thing to get right.

The localisation results carry a similar warning. The chain runs end to end on real hardware and reached its target in every configuration, but reprojection error did not predict metric accuracy, and calibration rather than vision dominated the residual budget. Both results say the same thing in different registers: the number a system reports on its own terms is a poor guide to how it will behave on someone else's.

Future work is to replace the target fiducial with a detected mask centroid, closing the two halves of this paper into a single markerless pipeline, and to add force feedback so the system moves from a metric target to a completed grasp.

## Data Availability

The annotation set (16,945 polygon masks), dataset audit tooling, evaluation code and trained detector weights are available at <https://github.com/Komil-jon/ev-battery-vision-pipeline>. The localisation pipeline, calibration procedure and UR5e validation scripts are available at <https://github.com/xcdgdj/D405-ArUco-UR5e-Validation>.

## Acknowledgements

The authors gratefully acknowledge the RESCu-M2 research hub at the University of Birmingham for supporting this work, and thank Prof. Samia Nefti-Meziani

for her leadership of the hub and for her guidance. We further thank the RESCu-M2 team for access to the EV battery disassembly facility and its imaging data, and to the robotics laboratory and its UR5e manipulator and RealSense camera, without which the localisation experiments could not have been carried out.

## References

- [1] European Commission, "Study on the critical raw materials for the EU: Final report," Publ. Office Eur. Union, Brussels, 2023.
- [2] G. Harper et al., "Recycling lithium-ion batteries from electric vehicles," *Nature*, vol. 575, pp. 75–86, 2019.
- [3] Transport & Environment, "From waste to value: the potential for battery recycling in Europe," Brussels, 2024.
- [4] European Parliament and Council, "Regulation (EU) 2023/1542 concerning batteries and waste batteries," *Off. J. Eur. Union*, vol. L191, pp. 1–117, 2023.
- [5] Y. Zang, M. Qu, D. T. Pham, R. Dixon, F. Goli, Y. Zhang, and Y. Wang, "Robotic disassembly of electric vehicle batteries: Technologies and opportunities," *Comput. Ind. Eng.*, vol. 198, 110727, 2024.
- [6] K. Wegener, S. Andrew, A. Raatz, K. Dröder, and C. Herrmann, "Disassembly of electric vehicle batteries using the example of the Audi Q5 hybrid system," *Procedia CIRP*, vol. 23, pp. 155–160, 2014.
- [7] M. Tan, J. Huang, X. Jiang, Y. Fang, Q. Liu, and D. T. Pham, "Robotic removal and collection of screws in collaborative disassembly of end-of-life electric vehicle batteries," *Biomimetics*, vol. 10, no. 8, 553, 2025.
- [8] "RAPID: Robotic agentic platform for intelligent electric vehicle disassembly," *arXiv:2603.18520*, 2026.
- [9] "Accurate screw detection method based on Faster R-CNN and rotation edge similarity for automatic screw disassembly," *Int. J. Comput. Integr. Manuf.*, vol. 34, no. 11, 2021.
- [10] "An accurate activate screw detection method for automatic electric vehicle battery disassembly," *Batteries*, vol. 9, no. 3, 187, 2023.
- [11] "Comparative analysis of YOLO and Faster R-CNN models for EV battery screw detection," *Int. J. Adv. Manuf. Technol.*, 2026.
- [12] "Computer vision application for industrial Li-ion battery module disassembly," *Prod. Eng.*, 2023.
- [13] S. Katiyar, P. Venigalla, K. Kosimov, and S. Nefti-Meziani, "Vision model for detection and condition assessment of EV battery components for circular manufacturing," in *Proc. 12th World Congr. Elect. Eng. Comput. Syst. Sci. (EECSS)*, London, UK, Aug. 2026, Paper MVML 125. doi: 10.11159/mvml26.125.
- [14] B. Recht, R. Roelofs, L. Schmidt, and V. Shankar, "Do ImageNet classifiers generalize to ImageNet?," in *Proc. ICML*, 2019, pp. 5389–5400.
- [15] P. W. Koh et al., "WILDS: A benchmark of in-the-wild distribution shifts," in *Proc. ICML*, 2021, pp. 5637–5664.
- [16] C. G. Northcutt, A. Athalye, and J. Mueller, "Pervasive label errors in test sets destabilize machine learning benchmarks," in *Proc. NeurIPS Datasets and Benchmarks*, 2021.
- [17] M. Oquab et al., "DINOv2: Learning robust visual features without supervision," *Trans. Mach. Learn. Res.*, 2024.
- [18] Roboflow, "RF-DETR: A real-time detection transformer," 2025.
- [19] G. Jocher and J. Qiu, "Ultralytics YOLO11," 2024.
- [20] G. Halmetschlager-Funek, M. Suchi, M. Kampel, and M. Vincze, "An empirical evaluation of ten depth cameras," *IEEE Robot. Autom. Mag.*, vol. 26, no. 1, pp. 67–77, 2019.
- [21] S. Garrido-Jurado, R. Muñoz-Salinas, F. J. Madrid-Cuevas, and M. J. Marín-Jiménez, "Automatic generation and detection of highly reliable fiducial markers under occlusion," *Pattern Recognit.*, vol. 47, no. 6, pp. 2280–2292, 2014.
- [22] V. Lepetit, F. Moreno-Noguer, and P. Fua, "EPnP: An accurate  $O(n)$  solution to the PnP problem," *Int. J. Comput. Vis.*, vol. 81, no. 2, pp. 155–166, 2009.